


✓ Q1: Time Series Analysis:



```
import pandas as pd
from pandas import read_csv

data = read_csv('https://raw.githubusercontent.com/ourcodingclub/CC-time-series/master/monthly_milk.csv',
                index_col='month',
                parse_dates = True)
```

✓ a. Find first five values of time series data (Monthly milk)

```
data.head()
```



| milk_prod_per_cow_kg | |  |
|----------------------|--------|---|
| month | |  |
| 1962-01-01 | 265.05 | |
| 1962-02-01 | 252.45 | |
| 1962-03-01 | 288.00 | |
| 1962-04-01 | 295.20 | |
| 1962-05-01 | 327.15 | |

Next steps:

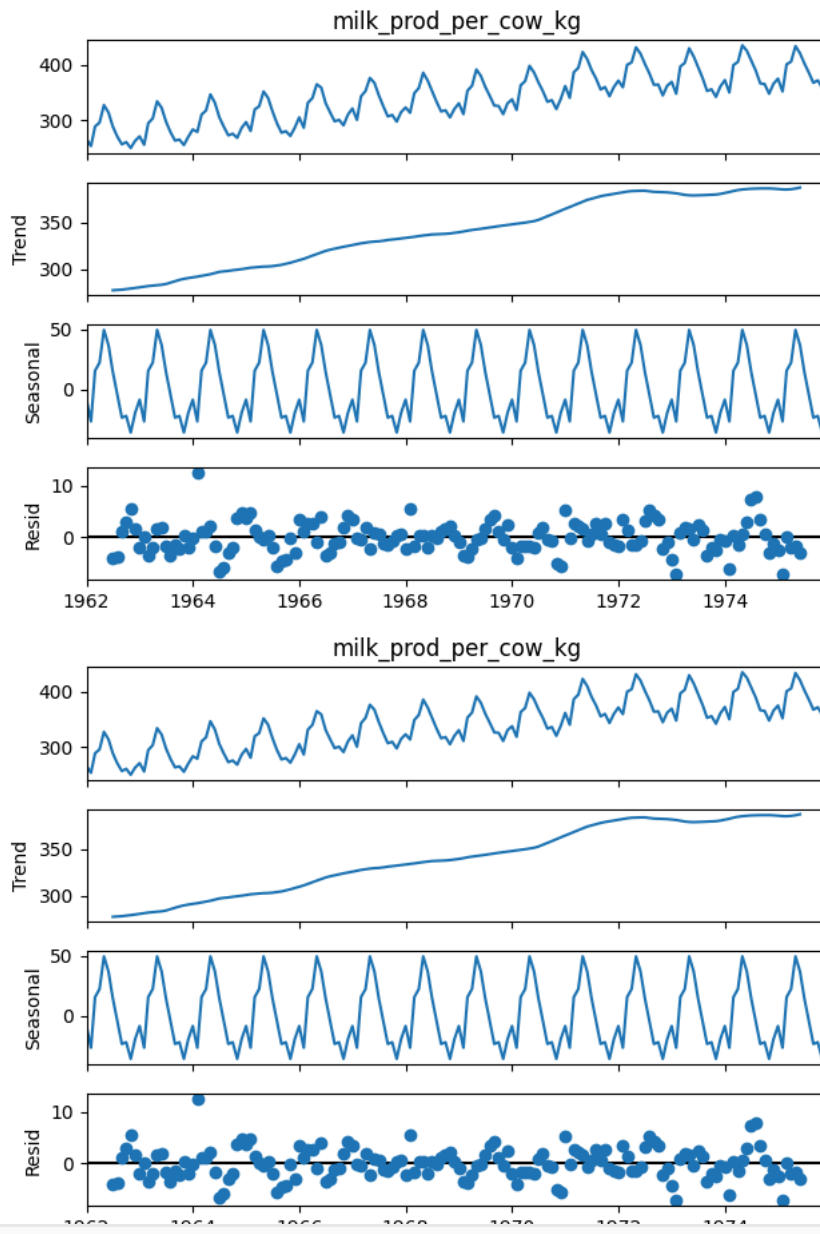
[Generate code with data](#)[View recommended plots](#)[New interactive sheet](#)

✓ b. Fit ARIMA model in the time series data set and predict for 2 years.

```
from statsmodels.tsa.seasonal import seasonal_decompose

# ETS Decomposition
result = seasonal_decompose(data['milk_prod_per_cow_kg'],
                             model='Additive')

# ETS plot
result.plot()
```



```
!pip install pmdarima
```



[Show hidden output](#)

```
import matplotlib.pyplot as plt
from pmdarima import auto_arima

import warnings
warnings.filterwarnings("ignore")

stepwise_fit = auto_arima(data['milk_prod_per_cow_kg'],
                           start_p = 1,
                           start_q = 1,
                           max_p = 3,
                           max_q = 3,
                           m = 12,
                           start_P = 0,
                           seasonal = True,
                           d = None,
```

```
D = 1,
trace = True,
error_action = 'ignore',
suppress_warnings = True,
stepwise = True)
```

```
stepwise_fit.summary()
```

```
➤ Performing stepwise search to minimize aic
ARIMA(1,0,1)(0,1,1)[12] intercept : AIC=823.289, Time=0.66 sec
ARIMA(0,0,0)(0,1,0)[12] intercept : AIC=1075.573, Time=0.05 sec
ARIMA(1,0,0)(1,1,0)[12] intercept : AIC=843.201, Time=0.38 sec
ARIMA(0,0,1)(0,1,1)[12] intercept : AIC=962.793, Time=0.38 sec
ARIMA(0,0,0)(0,1,0)[12] intercept : AIC=1203.299, Time=0.06 sec
ARIMA(1,0,1)(0,1,0)[12] intercept : AIC=870.000, Time=0.19 sec
ARIMA(1,0,1)(1,1,1)[12] intercept : AIC=825.170, Time=0.77 sec
ARIMA(1,0,1)(0,1,2)[12] intercept : AIC=825.147, Time=1.42 sec
ARIMA(1,0,1)(1,1,0)[12] intercept : AIC=839.342, Time=0.53 sec
ARIMA(1,0,1)(1,1,2)[12] intercept : AIC=inf, Time=7.51 sec
ARIMA(1,0,0)(0,1,1)[12] intercept : AIC=826.588, Time=0.38 sec
ARIMA(2,0,1)(0,1,1)[12] intercept : AIC=824.908, Time=0.80 sec
ARIMA(1,0,2)(0,1,1)[12] intercept : AIC=824.202, Time=0.71 sec
ARIMA(0,0,0)(0,1,1)[12] intercept : AIC=1075.819, Time=0.29 sec
ARIMA(0,0,2)(0,1,1)[12] intercept : AIC=918.974, Time=0.64 sec
ARIMA(2,0,0)(0,1,1)[12] intercept : AIC=822.921, Time=0.53 sec
ARIMA(2,0,0)(0,1,0)[12] intercept : AIC=869.597, Time=0.16 sec
ARIMA(2,0,0)(1,1,1)[12] intercept : AIC=824.788, Time=0.76 sec
ARIMA(2,0,0)(0,1,2)[12] intercept : AIC=824.764, Time=1.25 sec
ARIMA(2,0,0)(1,1,0)[12] intercept : AIC=838.597, Time=0.51 sec
ARIMA(2,0,0)(1,1,2)[12] intercept : AIC=inf, Time=7.31 sec
ARIMA(3,0,0)(0,1,1)[12] intercept : AIC=824.888, Time=0.66 sec
ARIMA(3,0,1)(0,1,1)[12] intercept : AIC=823.661, Time=1.42 sec
ARIMA(2,0,0)(0,1,1)[12] intercept : AIC=827.433, Time=0.50 sec
```

```
Best model: ARIMA(2,0,0)(0,1,1)[12] intercept
Total fit time: 27.898 seconds
```

SARIMAX Results

| | | | |
|-----------------------|----------------------------------|--------------------------|----------|
| Dep. Variable: | y | No. Observations: | 168 |
| Model: | SARIMAX(2, 0, 0)x(0, 1, [1], 12) | Log Likelihood | -406.460 |
| Date: | Fri, 15 Nov 2024 | AIC | 822.921 |
| Time: | 17:36:35 | BIC | 838.170 |
| Sample: | 01-01-1962 | HQIC | 829.114 |
| | - 12-01-1975 | | |

Covariance Type: opg

| | coef | std err | z | P> z | [0.025 | 0.975] |
|-------------------------|---------|-------------------|--------|-------|--------|--------|
| intercept | 0.7228 | 0.407 | 1.776 | 0.076 | -0.075 | 1.520 |
| ar.L1 | 0.7309 | 0.084 | 8.741 | 0.000 | 0.567 | 0.895 |
| ar.L2 | 0.1888 | 0.080 | 2.359 | 0.018 | 0.032 | 0.346 |
| ma.S.L12 | -0.6167 | 0.073 | -8.405 | 0.000 | -0.761 | -0.473 |
| sigma2 | 10.2694 | 0.995 | 10.318 | 0.000 | 8.319 | 12.220 |
| Ljung-Box (L1) (Q): | 0.00 | Jarque-Bera (JB): | 42.73 | | | |
| Prob(Q): | 0.97 | Prob(JB): | 0.00 | | | |
| Heteroskedasticity (H): | 0.83 | Skew: | 0.78 | | | |
| Prob(H) (two-sided): | 0.52 | Kurtosis: | 5.03 | | | |

Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients (complex-step)
```

✓ Using SARIMAX

```
from statsmodels.tsa.statespace.sarimax import SARIMAX
```

```
model_SARIMAX = SARIMAX(data['milk_prod_per_cow_kg'],
                          order = (2,0,0),
                          seasonal_order =(0, 1, 1, 12))
```

```
result = model_SARIMAX.fit()
result.summary()
```



SARIMAX Results

Dep. Variable: milk_prod_per_cow_kg **No. Observations:** 168
Model: SARIMAX(2, 0, 0)x(0, 1, [1], 12) **Log Likelihood** -409.717
Date: Fri, 15 Nov 2024 **AIC** 827.433
Time: 17:44:00 **BIC** 839.633
Sample: 01-01-1962 **HQIC** 832.388
- 12-01-1975

Covariance Type: opg

| | coef | std err | z | P> z | [0.025 | 0.975] |
|----------|---------|---------|--------|-------|--------|--------|
| ar.L1 | 0.7699 | 0.078 | 9.879 | 0.000 | 0.617 | 0.923 |
| ar.L2 | 0.2221 | 0.076 | 2.904 | 0.004 | 0.072 | 0.372 |
| ma.S.L12 | -0.6145 | 0.071 | -8.616 | 0.000 | -0.754 | -0.475 |
| sigma2 | 10.6326 | 1.030 | 10.321 | 0.000 | 8.613 | 12.652 |

Ljung-Box (L1) (Q): 0.02 **Jarque-Bera (JB):** 37.20
Prob(Q): 0.90 **Prob(JB):** 0.00
Heteroskedasticity (H): 0.80 **Skew:** 0.74
Prob(H) (two-sided): 0.43 **Kurtosis:** 4.87

Warnings:

```
train = data.iloc[:len(data)-24]
test = data.iloc[len(data)-24:]

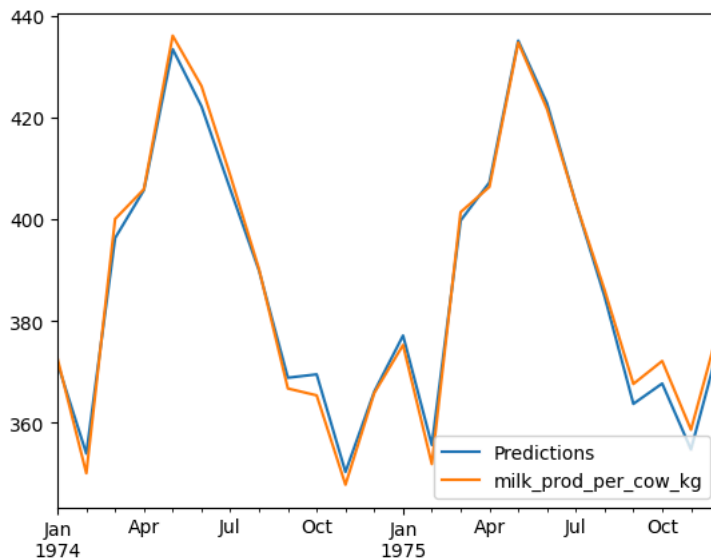
start = len(train)
end = len(train) + len(test) - 1

predictions = result.predict(start, end,
                             typ = 'levels').rename("Predictions")

predictions.plot(legend = True)
test['milk_prod_per_cow_kg'].plot(legend = True)
```



<Axes: xlabel='month'>



Using ARIMA

```
from statsmodels.tsa.arima.model import ARIMA
from pandas import DataFrame
from matplotlib import pyplot

model_ARIMA = ARIMA(data['milk_prod_per_cow_kg'], order=(2,0,0))
```

```

model_fit = model_ARIMA.fit()

predictions = model_fit.predict(start, end)
print(predictions)


pyplot.plot(data['milk_prod_per_cow_kg'], label='Original Data')
pyplot.plot(predictions, color='red', label='Predictions')
pyplot.legend()
pyplot.show()

```

```

↔ 1974-01-01    360.024701
   1974-02-01    369.946099
   1974-03-01    347.110642
   1974-04-01    397.533788
   1974-05-01    399.328535
   1974-06-01    428.245168
   1974-07-01    416.243120
   1974-08-01    399.918757
   1974-09-01    383.315574
   1974-10-01    361.960158
   1974-11-01    362.474123
   1974-12-01    345.481393
   1975-01-01    364.389984
   1975-02-01    372.189651
   1975-03-01    348.653247
   1975-04-01    398.708330
   1975-05-01    399.661421
   1975-06-01    426.894740
   1975-07-01    411.964482
   1975-08-01    395.009526
   1975-09-01    379.791948
   1975-10-01    363.153587
   1975-11-01    368.980023
   1975-12-01    355.475742
Freq: MS, Name: predicted_mean, dtype: float64

```



✓ Q2 : Linear Regression Problem:

```

df = pd.read_csv('mtcars.csv')

df.head()

```

| | model | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb |
|---|----------------|------|-----|-------|-----|------|-------|-------|----|----|------|------|
| 0 | Mazda RX4 | 21.0 | 6 | 160.0 | 110 | 3.90 | 2.620 | 16.46 | 0 | 1 | 4 | 4 |
| 1 | Mazda RX4 Wag | 21.0 | 6 | 160.0 | 110 | 3.90 | 2.875 | 17.02 | 0 | 1 | 4 | 4 |
| 2 | Datsun 710 | 22.8 | 4 | 108.0 | 93 | 3.85 | 2.320 | 18.61 | 1 | 1 | 4 | 1 |
| 3 | Hornet 4 Drive | 21.4 | 6 | 258.0 | 110 | 3.08 | 3.215 | 19.44 | 1 | 0 | 3 | 1 |

Next steps:

[Generate code with df](#)

[View recommended plots](#)


[New interactive sheet](#)

- Create a subset of the variables mpg, disp, hp and wt from the mtcars data. Construct the scatter plot for these attributes.

```
import matplotlib.pyplot as plt

subset = df[['mpg', 'disp', 'hp', 'wt']]

pd.plotting.scatter_matrix(subset, figsize=(10, 10), diagonal='kde')
plt.show()
```



- ✓ b. Fit a multiple linear regression model with mpg as the response variable and disp and hp as the input variables

```
import statsmodels.formula.api as sm

# Define the model formula
formula = 'mpg ~ disp + hp'

# Fit the model
model = sm.ols(formula=formula, data=df).fit()

print(model.summary())
```



OLS Regression Results

| | | | |
|-------------------|------------------|---------------------|----------|
| Dep. Variable: | mpg | R-squared: | 0.748 |
| Model: | OLS | Adj. R-squared: | 0.731 |
| Method: | Least Squares | F-statistic: | 43.09 |
| Date: | Fri, 15 Nov 2024 | Prob (F-statistic): | 2.06e-09 |
| Time: | 17:58:04 | Log-Likelihood: | -80.309 |
| No. Observations: | 32 | AIC: | 166.6 |
| Df Residuals: | 29 | BIC: | 171.0 |
| Df Model: | 2 | | |
| Covariance Type: | nonrobust | | |

| | coef | std err | t | P> t | [0.025 | 0.975] |
|-----------|---------|---------|--------|-------|--------|--------|
| Intercept | 30.7359 | 1.332 | 23.083 | 0.000 | 28.013 | 33.459 |
| disp | -0.0303 | 0.007 | -4.098 | 0.000 | -0.045 | -0.015 |
| hp | -0.0248 | 0.013 | -1.856 | 0.074 | -0.052 | 0.003 |

| | | | |
|----------------|-------|-------------------|-------|
| Omnibus: | 3.082 | Durbin-Watson: | 1.370 |
| Prob(Omnibus): | 0.214 | Jarque-Bera (JB): | 2.788 |
| Skew: | 0.680 | Prob(JB): | 0.248 |
| Kurtosis: | 2.508 | Cond. No. | 733. |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

- ✓ c. Check the p-values to evaluate the significance of the input variables.

```
# From the 'P>|t|' column in the output of model.summary().

# p value for 'disp': 0.000, less than 0.05. There's enough evidence to reject the null hypothesis.

# p value for 'hp': 0.074, greater than 0.05. There's not enough evidence to reject the null hypothesis.
```

... d. Which independent variable is the least significant? Why?